

Spring 2022 Introduction to Artificial Intelligence

Report of Homework #3

Student name: 劉子齊 Jonathan
Student ID: 0716304

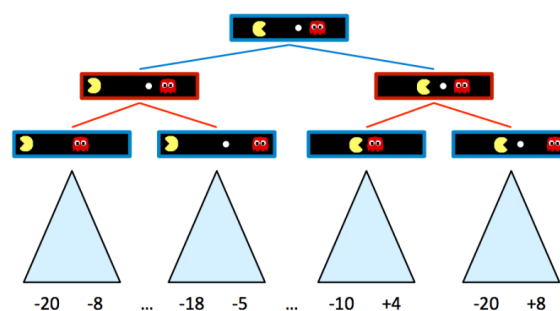
Part 1 : Adversarial search

Part 1-1: Minimax Search

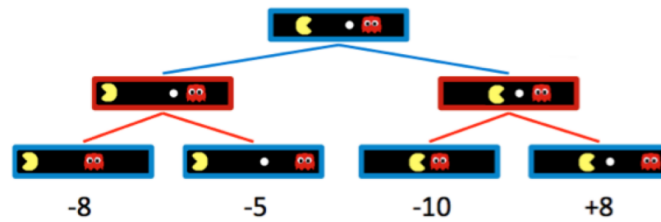
When the computer wants to make the next move in Pacman, of course it will play the one with the most points, which is eating all the food in this case, but this can easily fall into the opponent's trap, which is being eaten by the ghosts.

Therefore, a reasonable idea is to divide all levels into two categories of enemy and us, which will be Pacman and the Ghost in our case. The more points the level under our side has, the better, and the level under the other side loses as few points as possible.

Also, we can't assume that the opponent is a fool, so at every level, we have to think that "the opponent may make the move that will cost us the most points", and we must choose the strategy of "minimizing maximum points loss" as much as possible. As a result, this strategy makes the Minimax Search.



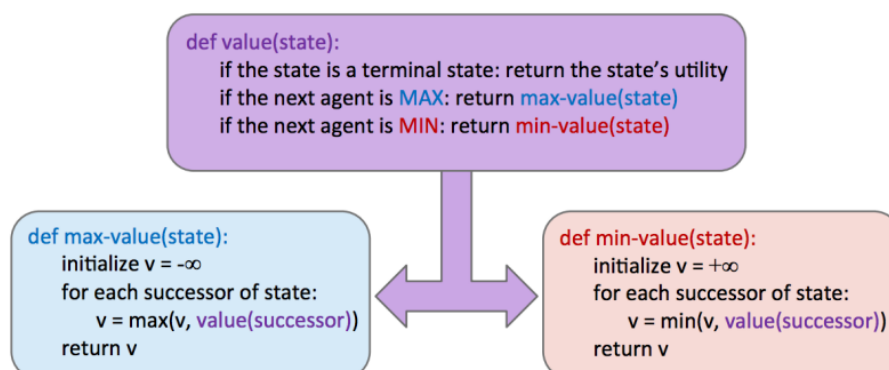
In the figure above, blue nodes correspond to nodes that Pacman controls and can decide what action to take, which will pick the choice with the max value, while red nodes correspond to ghost-controlled nodes, which will pick the minimum. Now let's move onto the second depth of the tree in the figure above.



As we can observe from the figure above, the blue Pacman nodes chose the option with the max value, which Pacman believes to be the best choice. The minimax algorithm only maximizes over the children of nodes controlled by Pacman, while minimizing over the children of nodes controlled by ghosts. Hence, the two ghost nodes above have values of $\min(-8, -5) = -8$ and $\min(-10, +8) = -10$ respectively. Correspondingly, the root node controlled by Pacman has a value of $\max(-8, -10) = -8$. As a result, the Pacman will get -8 as the score of this game.

However, if I was the Pacman in this game, I will not be satisfied, since I knew that I could have get $+8$ as my final score, if I chose the right way. Hence, we will need to have the Pacman put some bet to move to way which is not so straightforward and will take some risks correspondingly.

When implementing the Minimax algorithm, I found it kind of similar to DFS (Deep First Search), which they both start with the leftmost terminal node and all the way to the right of the game tree. Basically, I followed the pseudocode found on a website of UC Berkeley in the figure below.



```

104 class MinimaxAgent(MultiAgentSearchAgent):
105     """
106     Your minimax agent (par1-1)
107     """
108     def getAction(self, gameState):
109 >         """ ...
131         """ YOUR CODE HERE """
132         # Begin your code
133
134         pacman_legal_actions = gameState.getLegalActions(0)
135
136         # print()
137         # print(pacman_legal_actions)
138
139         max_value = float('-inf')
140         decision = None
141
142         for action in pacman_legal_actions:
143
144             action_score = self.Min_Value(gameState.getNextState(0, action), 1, 0)
145             # print("action: ", action, "action_score: ", action_score)
146
147             if ((action_score) > max_value ):
148                 max_value = action_score
149                 decision = action
150
151         # print("final_call", decision)
152         return decision
153
154         util.raiseNotDefined()
155         # End your code
156

```

At the very first of my program in the figure above, I call the `getLegalActions()` function to get all the possible actions the Pacman can move for the next step. Besides, I declared the initial `max_value`, which determines the final action, as a negative infinite float number, and the decision as a NULL.

After declaring, I construct a loop which depends on the elements in the list named “`pacman_legal_actions`”. For each action in the list, I first obtain the “`action_score`” by the Minimax algorithm, then I compare all of them throughout the whole loop and get the final decision, which is the action with the highest “`action_score`”.

```

157     def Max_Value (self, gameState, depth):
158
159         # end of the game
160         if(depth == self.depth):
161             return self.evaluationFunction(gameState)
162         # dead
163         elif(len(gameState.getLegalActions(0)) == 0):
164             return self.evaluationFunction(gameState)
165
166         # print("Max_Value Self Depth: ", self.depth)
167
168         return max([self.Min_Value(gameState.getNextState(0, action), 1, depth)
169                     for action in gameState.getLegalActions(0)])
170

```

For the Minimax algorithm, I separated it into two parts, one is for the minimum value, and the other for the maximum. For the `Max_Value()` part, I first confirm if the current depth equals to the “`self.depth`”, which is the end of the game tree. If we’ve reach the end of the game tree, then I call the `evaluationFunction()` to get the final result and return it. Also, I’ll check whether there are no legal actions left, which might mean the Pacman just ran into a ghost and died. In this case, I’ll also return the final result obtained by the `evaluationFunction()`.

If none of the two scenarios mentioned above happened, I’ll call the `Min_Value` function, which will be mentioned later, and iterate it through all the legal actions obtained by the `getLegalActions(0)`, which returns a list contains all the legal actions of Pacman since “0” stands for Pacman here. As I keep getting value from the `Min_Value` function, I’ll compare them all and keep the maximum of all of them. Finally, I’ll return the maximum value.

```
172     def Min_Value (self, gameState, agentIndex, depth):
173
174         # dead
175         if(len(gameState.getLegalActions(agentIndex)) == 0):
176             return self.evaluationFunction(gameState)
177
178         # when the game continues
179         if(agentIndex < gameState.getNumAgents() - 1):
180             return min([self.Min_Value(gameState.getNextState(agentIndex, action), agentIndex + 1, depth)
181                         for action in gameState.getLegalActions(agentIndex)])
182         # when the last ghost remaining
183         else:
184             return min([self.Max_Value(gameState.getNextState(agentIndex, action), depth + 1)
185                         for action in gameState.getLegalActions(agentIndex)])
```

For the `Min_Value()` part, I’ll check whether there are no legal actions left, which might means the Pacman just ran into a ghost and died. In this case, I’ll also return the final result of the current `gameState` obtained by the `evaluationFunction()`.

Then I’ll check if there’s still agents remaining. If there are still agents remaining, I’ll call the `Min_Value` function, and iterate it through all the legal actions obtained by the `getLegalActions(agentIndex)`, which returns a list contains all the legal actions of the corresponding agent. As I keep getting value from the `Min_Value` function, I’ll sum them up and return the final value.

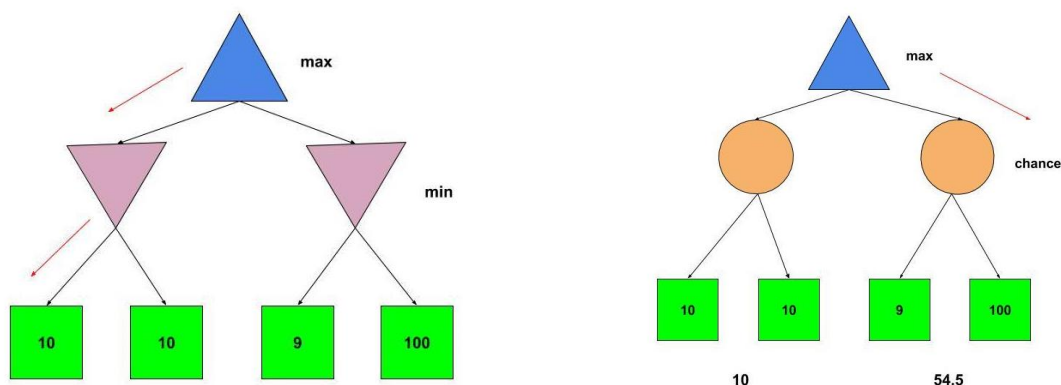
If there’s no agents remaining, I’ll call the `Max_Value` function, and iterate it through all the depth and all the legal actions obtained by the

getLegalActions(agentIndex), which returns a list contains all the legal actions of the corresponding agent. As I keep getting value from the Min_Value function, I'll compare them all and keep the minimum of all of them. Finally, I'll return the minimum value.

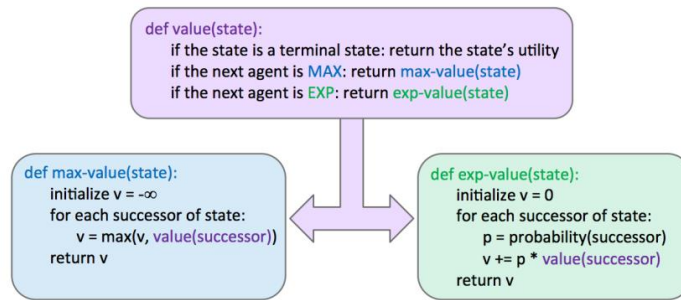
Part 1-2: Expectimax Search

I'll take Expectimax as a variation of Minimax Algorithm, which replaces minimizer nodes into chance nodes in the original Minimax game tree. As we know that the minimizer plays optimally, it makes sense to go to the left. But what if there is a possibility of the minimizer making a mistake. Therefore, going right might sound more appealing or may result in a better solution.

Hence, in the Expectimax Search, I replace minimizer nodes by chance nodes, which takes the average of the agent numbers, which can be seen in the figure below. On the left, we have the original Minimax Search, and on the left we have the tree of Expectimax Search.



For the implementation, I followed the pseudocode found on a website of UC Berkeley in the figure below. Which we can find both of them are basically the same, except for changing the Min_Value() function of the Minimax Search into Exp_Value() function which contains the chance node, of the Expectimax Search.



```

187 class ExpectimaxAgent(MultiAgentSearchAgent):
188     """
189     Your expectimax agent (part1-2)
190     """
191
192     def getAction(self, gameState):
193         """
194         Returns the expectimax action using self.depth and self.evaluationFunction
195
196         All ghosts should be modeled as choosing uniformly at random from their
197         legal moves.
198         """
199         """ YOUR CODE HERE """
200         # Begin your code
201
202         pacman_legal_actions = gameState.getLegalActions(0)
203
204         # print()
205         # print(pacman_legal_actions)
206
207         max_value = float('-inf')
208         decision = None
209
210         for action in pacman_legal_actions:
211
212             action_score = self.Exp_Value(gameState.getNextState(0, action), 1, 0)
213             # print("action: ", action, "action_score: ", action_score)
214
215             if ((action_score) > max_value ):
216                 max_value = action_score
217                 decision = action
218
219         # print("final_call", decision)
220         return decision
221
222         util.raiseNotDefined()
223         # End your code

```

```

225 def Max_Value (self, gameState, depth):
226
227     # end of the game
228     if(depth == self.depth):
229         return self.evaluationFunction(gameState)
230     # dead
231     elif(len(gameState.getLegalActions(0)) == 0):
232         return self.evaluationFunction(gameState)
233
234     # print("Max_Value Self Depth: ", self.depth)
235
236     return max([self.Exp_Value(gameState.getNextState(0, action), 1, depth)
237                 for action in gameState.getLegalActions(0)])

```

By the figure, we can find the `get_action()` function and the `Max_Value()` function are almost the same. I also checked the same condition and return the

same value in both functions as I did in the Minimax Search. The only thing changed was in the Max_Value() function, which I changed the original Min_Value() function into Exp_Value() function.

```
240 def Exp_Value (self, gameState, agentIndex, depth):
241
242     num_actions = len(gameState.getLegalActions(agentIndex))
243
244     # dead
245     if(len(gameState.getLegalActions(agentIndex)) == 0):
246         return self.evaluationFunction(gameState)
247
248     # when the game continues
249     if(agentIndex < gameState.getNumAgents() - 1):
250         return sum([self.Exp_Value(gameState.getNextState(agentIndex, action), agentIndex + 1, depth)
251                     for action in gameState.getLegalActions(agentIndex)]) / float(num_actions)\
252     # when the last ghost remaining
253     else:
254         return sum([self.Max_Value(gameState.getNextState(agentIndex, action), depth + 1)
255                     for action in gameState.getLegalActions(agentIndex)]) / float(num_actions)
```

In the Exp_Value() function, basically, it's also kind of similar to the Min_Vlaue() function of the Minimax Search. I checked whether there are no legal actions left, which might mean the Pacman just ran into a ghost and died. In this case, I'll also return the final result of the current gameState obtained by the evaluationFunction().

Then I'll check if there's still agents remaining. If there are still agents remaining, I'll call the Exp_Value() function, and iterate it through all the legal actions obtained by the getLegalActions(agentIndex), which returns a list contains all the legal actions of the corresponding agent. As I keep getting value from the Exp_Value() function, I'll sum them up. After getting the final sum value, I divided the sum with the total num of the actions and return the final value.

If there's no agents remaining, I'll call the Max_Value() function, and iterate it through all the depth and all the legal actions obtained by the getLegalActions(agentIndex), which returns a list contains all the legal actions of the corresponding agent. As I keep getting value from the Max_Value() function, I'll sum them up. After getting the final sum value, I divided the sum with the total num of the actions. Finally, I'll return the final value.

Part 1-3: Evaluation Function (Bonus)

```
258 def betterEvaluationFunction(currentGameState):
259     """ ...
260     """
261     """ *** YOUR CODE HERE *** """
262     # Begin your code
263
264     pacman_position = currentGameState.getPacmanPosition()
265     ghost_positions = currentGameState.getGhostPositions()
266
267     #number of remaining capsule
268     capsule_num = len(currentGameState.getCapsules())
269
270     # print()
271     # print(currentGameState.getCapsules())
272
273     game_score = currentGameState.getScore()
274
275     # ----- food area -----
276     # get the position of all food
277     food_place = currentGameState.getFood()
278     food_place = food_place.asList()
279
280     #number of remaining food
281     food_num = len(food_place)
282
283     # print()
284     # print(food_place)
285
286     closest_food = 1
287     food_distances = [manhattanDistance(pacman_position, food_position) for food_position in food_place]
288     # print("food_distances: ", food_distances)
289
290     if food_num > 0:
291         closest_food = min(food_distances)
292
293     #distance of ghost to us
294     for ghost_position in ghost_positions:
295         ghost_distance = manhattanDistance(pacman_position, ghost_position)
296
297         # If ghost is too close to us, then its time to stop eating and escape lol
298         if ghost_distance < 5:
299             closest_food = 878787
300
301     #calculate all the weighted features and sum them up
302     features = [(1.0 / closest_food), game_score, food_num, capsule_num]
303     weights = [10, 200, -100, -10]
304     return sum([feature * weight for feature, weight in zip(features, weights)])
305
306     util.raiseNotDefined()
307     # End your code
```

In part 1-3, I redefined the way the program evaluates the better option of the following move. First, I declared the two variables, “pacman_position” and “ghost_positions”, which obtains the current position of the Pacman and all the ghosts in the current match. Then I get the current “food_num” and the “capsule_num” to know how many food or capsules are left in the map. Also, I declared “food_place” which is a list that store the position of all the food in the form of a two-dimensional coordinate.

After all the variables are set, I set the value of “closest_food” to 1, which “closest_food” indicated the distance of the closest food to the Pacman. Then I

get a list of all the food distances by iterating through all the food positions in the “food_place” list and calculate the distance through Manhattan distance. With the list of distances to all the food, if there’s still food left in the map, I’ll set the value of “closest_food” to the minimum value in the list of food distances.

Next, I started to deal with the ghosts. For every ghost position in the list named “ghost_positions”, I obtain the distance from the Pacman to the ghost through Manhattan distance. When the distance is smaller than 5, my Pacman will leave the food alone, and start to escape to keep itself alive.

For last, I defined a list containing the reciprocal of “closest_food”, “game_score”, “food_num”, and “capsule_num”. Also, the other list with all the corresponding weight that I would like to give every variable in the previous list, which are 10, 200, -100, and -10. Then I iterate through them simultaneously through the zip() function and multiply them. Finally, I sum the result of multiplication up, and return it as the final result.

Part 2: Q-learning

Part 2-1: Value Iteration

In this part batch version of Value Iteration is implemented. The values of states in previous iteration is used for updating values of states in next iteration which might not always be the case with online update.

The `getStates()` function returns all the states in the environment. The function `getPossibleActions(state)` returns the possible actions from state. The function `getQValue(state, action)` returns the QValue. I iterated through all the iterations, states and actions one by one, then get the corresponding value through the three functions mentioned above.

After calculating every state value through the functions mentioned above, I updated the value for each state by using the `argmax()` function. Finally, I copy back the new value through the `copy()` function and return it.

```
67     def runValueIteration(self):
68         # Write value iteration code here
69         """ YOUR CODE HERE """
70         # Begin your code
71
72         for state in self.mdp.getStates():
73             self.values[state] = float(0)
74
75         for iteration in range(self.iterations):
76             next_values = self.values.copy()
77
78             for state in self.mdp.getStates():
79                 state_values = util.Counter()
80
81                 for action in self.mdp.getPossibleActions(state):
82                     state_values[action] = self.getQValue(state, action)
83
84                 next_values[state] = state_values[state_values.argmax()]
85
86             self.values = next_values.copy()
87
88         return self.values
89
90         util.raiseNotDefined()
91         # End your code
```

In the `computeQValueFromValues()` function in the figure below, I first obtain all the transition probabilities through the `getTransitionsStatesAndProbs(state, action)` function. Also, I set the initial value of the QValue to 0.0, which is a float type number.

For every transition probability in all the transition probabilities obtained, I extract the Tstate and the probability in each transition. With the Tstate value and the probability, I updated the Qvalue through summing up every possible value in all the transitions. By following the function below, I obtain all the QValues in this

function.

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Note: I only calculate all the Qvalues in this function, I'll determine the max Qvalue and return it in the next function.

```
101 def computeQValueFromValues(self, state, action):
102     """
103     Compute the Q-value of action in state from the
104     value function stored in self.values.
105     """
106     """ YOUR CODE HERE """
107     # Begin your code
108
109     transition_probs = self.mdp.getTransitionStatesAndProbs(state, action)
110     QValue = float(0)
111
112     for transition in transition_probs:
113         Tstate, prob = transition
114         QValue += prob * (self.mdp.getReward(state, action, Tstate) + self.discount * self.getValue(Tstate))
115
116     return QValue
117
118     util.raiseNotDefined()
119     # End your code
```

In the computeActionFromValue() function, I first check whether the game is about to terminal, if not, I will start to calculate.

First, I initiate the QValue by the Counter() function in the utisl.py, which is a dictionary with a default value of zero. Also, I obtained all the actions through the getPossibleActions().

With all the actions, I get the QValue for each action in a list, then return the maximum QValue with the argMax() function, just exactly as the following function I mentioned earlier.

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

```
121 def computeActionFromValues(self, state):
122     > """ ...
130     """ YOUR CODE HERE """
131     # Begin your code
132
133     #check for terminal
134
135     if(self.mdp.isTerminal(state)):
136         return None
137     else:
138         QValues = util.Counter()
139         actions = self.mdp.getPossibleActions(state)
140
141         for action in actions:
142             QValues[action] = self.computeQValueFromValues(state, action)
143
144         return QValues.argmax()
145
146     util.raiseNotDefined()
```

Part 2-2: Q-learning

In the very first `__init__()` function, the only thing I did is I initiated the Qvalues to a dictionary with a default value of zero by the `Counter()` function in the `util.py`.

```
45     def __init__(self, **args):
46         "You can initialize Q-values here..."
47         ReinforcementAgent.__init__(self, **args)
48
49         """ YOUR CODE HERE """
50         # Begin your code
51
52         self.QValues = util.Counter()
53         print("Alpha:", self.alpha)
54         print("Discount: ", self.discount)
55         print("Exploration: ", self.epsilon)
56
57         # End your code
58
```

In the `getQValue()` function, I return the Qvalues, which should return the Q Value if the state actually exists, and will return 0.0, as I initiated in the `__init__()` function if the state didn't exist.

```
60     def getQValue(self, state, action):
61         """
62         Returns Q(state,action)
63         Should return 0.0 if we have never seen a state
64         or the Q node value otherwise
65         """
66         """ YOUR CODE HERE """
67         # Begin your code
68
69         return self.QValues[(state, action)]
70
71         util.raiseNotDefined()
72         # End your code
73
```

In the `computeValueFromQvalues()` function, I first obtain all the Q values for all the legal actions through the `getQValue()` function. If I didn't get any value, which means there's no legal actions, I'll return 0.0 for the result. Else, I'll return the maximum element from all the values.

```
75     def computeValueFromQvalues(self, state):
76         """ ...
77
78         """ YOUR CODE HERE """
79         # Begin your code
80
81         values = [self.getQValue(state, action) for action in self.getLegalActions(state)]
82
83         if(values):
84             return max(values)
85         else:
86             return 0.0
87
88         util.raiseNotDefined()
89         # End your code
90
```

For the `computeActionFromQvalues()` function, for each action from all the legal actions, I calculate its Q value through the `getQValue()` function. If the QValue is greater than the `max_QValue` or there's no max action, then I'll update the `max_QValue` and the `max_action` through the Qvalue and the action correspondingly. After iterating through all the legal actions, I will return the final biggest `max_action` as the final result.

```

95 def computeActionFromQValues(self, state):
96     """ ...
101     """ *** YOUR CODE HERE ***
102     # Begin your code
103
104     max_action = None
105     max_Qvalue = 0
106
107     for action in self.getLegalActions(state):
108         Qvalue = self.getQValue(state, action)
109
110         if(Qvalue > max_Qvalue or max_action == None):
111             max_Qvalue = Qvalue
112             max_action = action
113
114     return max_action
115
116     util.raiseNotDefined()
117     # End your code

```

For the `update()` function, I divided it into two parts. For the first part, I multiplied the Q value obtained by the `getQValue()` function with $(1 - \alpha)$. For the second part, I first check if there still exist the next action, if not, the sample will be set to the value of the reward. Else, the sample value will be the reward plus the maximum value obtained from the `getQValue` function through all the legal next actions. Then multiplies the sample value with α , which results the second part. Finally, I set the Q value of the current state and action into the addition of the first part and the second part.

```

147 def update(self, state, action, nextState, reward):
148     """ ...
149
150     """
151     """ *** YOUR CODE HERE *** """
152     # Begin your code
153
154     first_part = (1 - self.alpha) * self.getQValue(state, action)
155
156     if len(self.getLegalActions(nextState)) == 0:
157         sample = reward
158     else:
159         sample = reward + (self.discount * max([self.getQValue(nextState, next_action)
160                                                  for next_action in self.getLegalActions(nextState)]))
161
162     second_part = self.alpha * sample
163
164     self.QValues[(state, action)] = first_part + second_part
165
166     # util.raiseNotDefined()
167     # End your code

```

Part 2-3: epsilon-greedy action selection

For the Epsilon-Greedy Action Selection, I modify the `getAction()` function directly. The main idea here is that exploration of state space has to be done first and after getting good policy, exploitation of that policy has to be done. With epsilon probability choose a random action from a state otherwise choose the present optimal policy.

Hence, from my code, we can find that as the recommendation of the comments, I chose to use the `flipCoin()` function to decide whether to get the action from the `choice()` function to pick randomly from the list of legal actions, or from the `getPolicy()` function directly. After that, I'll return the action which I get from my final randomly decision.

```
118
119 ✓ def getAction(self, state):
120 >     """...
130     # Pick Action
131     legalActions = self.getLegalActions(state)
132     action = None
133     """ YOUR CODE HERE """
134     # Begin your code
135
136 ✓     if(util.flipCoin(self.epsilon)):
137         action = random.choice(legalActions)
138 ✓     else:
139         action = self.getPolicy(state)
140
141     return action
142
143     util.raiseNotDefined()
144     # End your code
145
```

Part 2-4: Approximate Q-learning

For approximate Q learning, I implemented through the function provided by the CS188 course of the UC Berkeley.

$$\begin{aligned}\text{difference} &= \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a) \\ Q(s, a) &\leftarrow Q(s, a) + \alpha [\text{difference}] \\ w_i &\leftarrow w_i + \alpha [\text{difference}] f_i(s, a)\end{aligned}$$

First of all, I implemented the `getQValue()` function. I obtained all the features by the `getFeatures()` function. For each feature obtained, I multiplied its weight, which I obtained it through the `getWeight()` function that can extract the weight of a particular function from the `getWeights()` function, with the feature. Then I sum them up as the “QValue” and return the final QValue.

Then I implement the `update()` function, which I basically followed the formula

mentioned above. I initiate the difference as “the reward” plus “the multiplication of the discount and the value of next state obtained by the `getValue()` function” subtract “the Q Value of the current state and action”. For each feature obtained by the `getFeatures` function, I set the weights for each feature as the original weight of the feature plus the multiplication of alpha, difference, and the feature value.

```

217 class ApproximateQAgent(PacmanQAgent):
218     """ ...
219
220     def __init__(self, extractor='IdentityExtractor', **args):
221         self.feateExtractor = util.lookup(extractor, globals())()
222         PacmanQAgent.__init__(self, **args)
223         self.weights = util.Counter()
224
225     def getWeights(self):
226         return self.weights
227
228     def getWeight(self, feature):
229         return self.weights[feature]
230
231     def getQValue(self, state, action):
232         """ ...
233
234         """ YOUR CODE HERE """
235         # Begin your code
236         # get weights and feature
237
238         features = self.feateExtractor.getFeatures(state, action)
239         QValue = 0.0
240
241         for feature in features :
242             QValue += self.getWeight(feature) * features[feature]
243
244         return QValue
245
246         util.raiseNotDefined()
247         # End your code
248
249     def update(self, state, action, nextState, reward):
250         """ ...
251
252         """ YOUR CODE HERE """
253         # Begin your code
254
255         features = self.feateExtractor.getFeatures(state, action)
256         difference = reward + (self.discount*self.getValue(nextState)) - self.getQValue(state, action)
257
258         for feature in features :
259             self.weights[feature] = self.getWeight(feature) + (self.alpha * difference * features[feature])
260
261         # util.raiseNotDefined()
262         # End your code
263
264     def final(self, state):
265         "Called at the end of each game."
266         # call the super-class final method
267         PacmanQAgent.final(self, state)
268 
```

Part 3: DQN

```

C:\Users\jonat\Desktop\Intro. to A.I\2022Spring_Artificial_Intelligence\hw3\DQN>
C:\Users\jonat\Desktop\Intro. to A.I\2022Spring_Artificial_Intelligence\hw3\DQN>python pacman.py -p PacmanDQN -n 10000 -x 10000 -l smallClassic
C:\Users\jonat\Desktop\Intro. to A.I\2022Spring_Artificial_Intelligence\hw3\DQN>
C:\Users\jonat\Desktop\Intro. to A.I\2022Spring_Artificial_Intelligence\hw3\DQN>python pacman.py -p PacmanDQN -n 200 -x 100 -l smallClassic
Started Pacman DQN algorithm
Model has been trained

```

